

## Lab 9: Indexes

### Part 1 - Selecting Indexes

**A database contains the following table for former-employee records**

The composite index on (“Start Date”, “End Date”) will efficiently support these queries. The “Start Date” part of the index will help find all employees who started on a certain date, and the “End Date” part of the index will assist in identifying those employees whose employment period includes the specified end date. Creating a single composite index saves us from adding separate indexes for each date column, reducing index maintenance overhead.

Summary of Additional Indexes:

- Index(“Start Date”, “End Date”)

**A database contains the following table for tracking student grades in classes**

1. Get all students with a grade better than ‘B’:
  - Additional Index: (Grade)
2. Get all classes where any student earned a grade worse than ‘D’:
  - Additional Index: (Grade)
3. Get all classes ordered by class name:
  - Additional Index: (className)
4. Get all students who earned an ‘A’ in a certain class:
  - Additional Index: (className, Grade)

After merging additional index from 3 & 4, we get Index: (className, Grade)

Summary of Additional Indexes:

- (Grade)
- (className, Grade)

### Queries on the chess database

1. `select Name from Players where Elo >= 2050;`
  - Additional Index: `Players.Elo`
2. `select Name, gID from Players join Games where pID = WhitePlayer;`
  - Existing Index: `Players.pID` (as it is the default primary key of Players table and already indexed).
  - Additional Index: `Games.WhitePlayer`

Summary of Additional Indexes:

- `Players.Elo`
- `Games.WhitePlayer`

## Queries on the public Library database

Analysis: - The **Inventory** table has a primary key (PK) on the **Serial** column.  
- The **CheckedOut** table also has a primary key (PK) on the **Serial** column. -  
Since **Serial** is the only shared column between **Inventory** and **CheckedOut**,  
and both tables have their primary keys indexed by default, there is no need to  
add any additional indexes for this specific query.

Summary:

- None

## More library queries

1. `select * from Inventory natural join CheckedOut where CardNum = 2;`
  - Natural join between **Inventory** and **CheckedOut** depends on the common primary key **Serial**.
  - Additional Index: **CheckedOut.CardNum**
2. `select * from Patrons natural join CheckedOut;`
  - **Patrons** PK is **CardNum**.
  - **CheckedOut** PK is **Serial**.
  - **Patrons** and **CheckedOut** share the common column **CardNum**.
  - **Patrons.CardNum** is already the primary key of **Patrons** table, so it is already indexed.

Summary of Additional Indexes:

- **CheckedOut.CardNum**

## Still more library queries

`Serials = from i in t.Inventory joins Titles with Inventory on Titles.ISBN = Inventory.ISBN.`

**Titles** has a primary key on **ISBN**, so it is already indexed. **Inventory** has a primary key on **Serial**, so we still need to index on **Inventory.ISBN**.

Summary of Additional Indexes:

- **Inventory.ISBN**

## Part 2 - B+ Tree Index Structures

**How many rows of the table can be placed into the first leaf node of the primary index before it will split?**

The primary index is created in the order (**studentID**, **className**). Let's calculate the size of each row in the index:

- **studentID** is an int, which occupies 4 bytes.

- `className` is a `varchar(10)`, which occupies 10 bytes.
- `Grade` is a `char(1)`, which occupies 1 byte.

Each row in the index occupies 15 bytes. The leaf node of the primary index has a size of 4096 bytes. Therefore, the number of rows that can be placed into the first leaf node of the primary index before it will split is  $\text{floor}(4096 / 15) = 273$ .

**What is the maximum number of keys stored in an internal node of the primary index? (Remember to ignore pointer space. Remember that internal nodes have a different structure than leaf nodes.)**

Available space in the internal node = 4096 bytes

Size of each key =  $4 + 10 = 14$  bytes

Maximum number of keys =  $\text{floor}(4096 \text{ bytes} / 14 \text{ bytes}) = 292$

**What is the maximum number of rows in the table if the primary index has a height of 1? (A tree of height 1 has 2 levels and requires exactly one internal node)**

From the last question we know that the B+ tree's order is  $292 + 1 = 293$ . And the max keys per internal node could hold is 292 keys.

The max leaf node number is  $293 \text{ leaf nodes} * 273 \text{ rows per leaf node} = 79989$  rows.

**What is the minimum number of rows in the table if the primary index has a height of 1? (A tree of height 1 has 2 levels). The minimum capacity of a node in a B+ tree is 50%, unless it is the only internal/leaf node. The minimum number of children of a root node is 2.**

The minimum number of leaf nodes is 2. The minimum number of rows per leaf node is  $\text{ceil}(273/2) = 137$ . Thus the minimum number of rows in the table is  $2 * 137 = 274$ .

**If there is a secondary index on Grade, what is the maximum number of entries a leaf node can hold in the secondary index?**

The Grade serves as the primary key, takes 1 byte. The pointer to the record takes 4 bytes. The total size of each entry is 5 bytes. The leaf node has a size of 4096 bytes. Therefore, the maximum number of entries a leaf node can hold in the secondary index is  $\text{floor}(4096 / 5) = 819$ .

**Another table**

**What is the maximum number of leaf nodes in the primary index if the table contains 48 rows?**

A leaf node could hold maximum  $\text{floor}(4096/128) = 32$  rows. A leaf node should at least hold  $\text{ceil}(32/2) = 16$  rows. The maximum number of leaf nodes is  $\text{ceil}(48/16) = 3$ .

**What is the minimum number of leaf nodes in the primary index if the table contains 48 rows?**

The minimum number of leaf nodes is  $\text{ceil}(48/32) = 2$ .